

# Convergence of Plug-and-Play algorithms for image inverse problems

**Samuel Hurault**  
ENS Paris

*PhD supervisors:* Arthur Leclaire, Nicolas Papadakis

*Collaborations:* Antonin Chambolle, Ulugbek Kamilov, Lingxiao Li, Justin Solomon,  
Matthieu Terris, Julian Tachella



# Image Inverse Problems



Acquisition



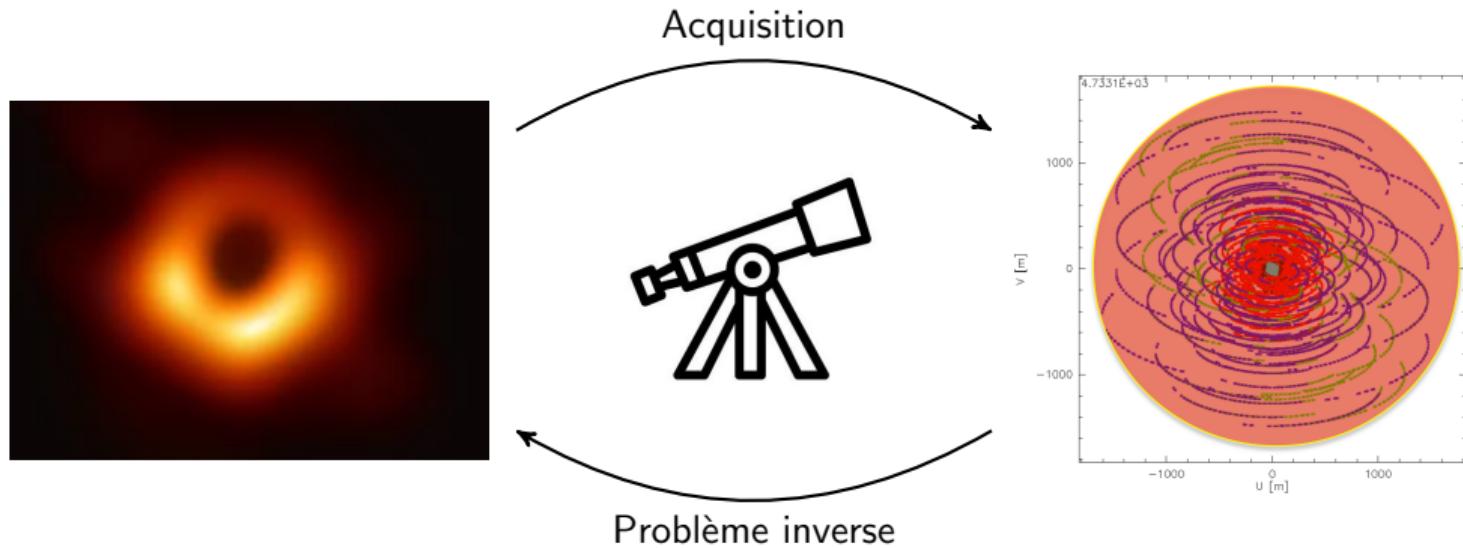
# Image Inverse Problems



# Image Inverse Problems

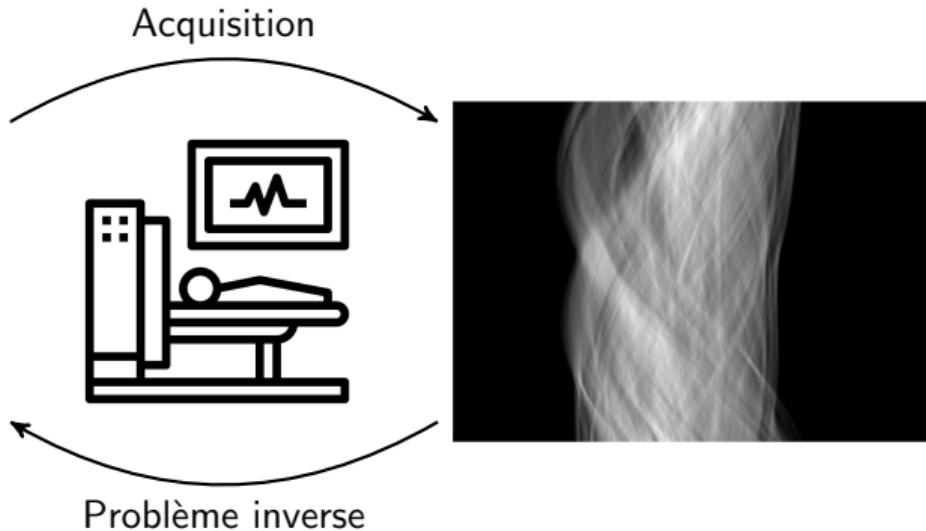
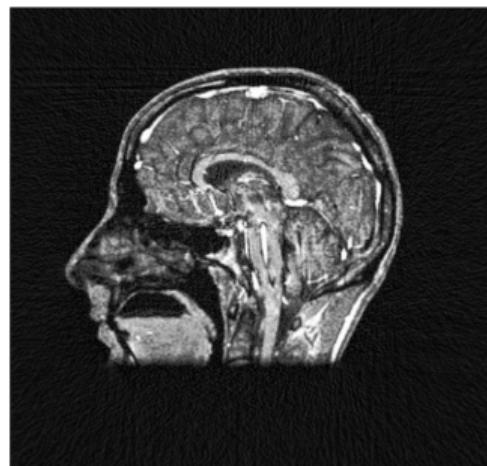


# Image Inverse Problems



Sources: Wikipedia, [\[Isella, '11\]](#)

# Image Inverse Problems

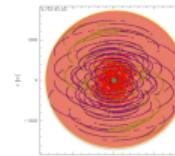
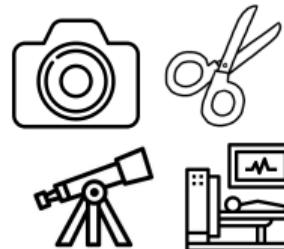


Source: Wikipedia

# Image Inverse Problems



Acquisition / Degradation

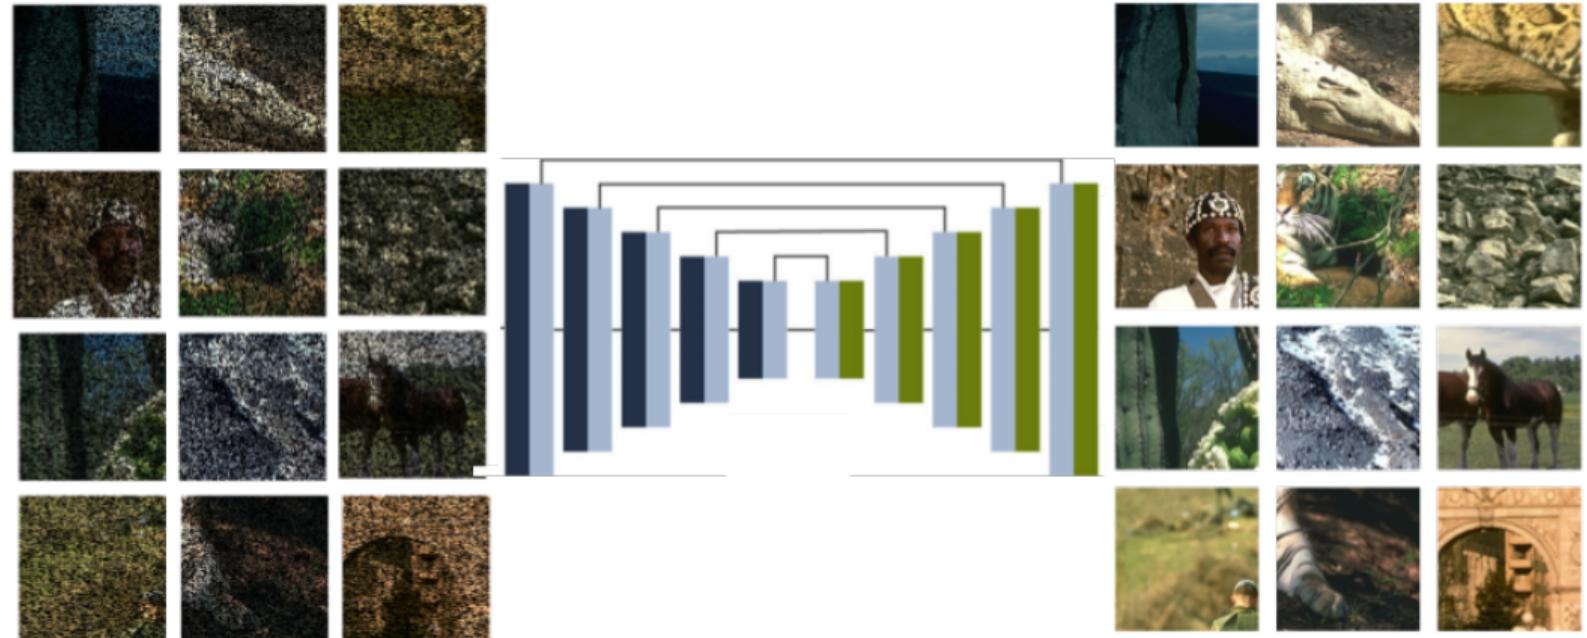


$$x \in \mathbb{R}^n$$

$$y \sim \text{noise}(Ax) \in \mathbb{R}^m$$

Inverse Problem

## Supervised methods [Alder & Ökten, '18, Bertocchi et. al '18, Zhang et. al '20]



- ✓ If access to a dataset  $(x_i, y_i)$ , state-of-the-art performance.
- ✗ Requires to retrain for each degradation  $A$ .

## Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

# Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

# Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

# Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

Data-fidelity

ex:  $f(x) = \frac{1}{2} \|Ax - y\|^2$

# Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

Data-fidelity

ex:  $f(x) = \frac{1}{2} \|Ax - y\|^2$

Regularization

ex: Total Variation [Rudin et al. '92]

Wavelet sparsity [Mallat '09]

Deep prior [Bora et al. '17]

## Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

## Probabilistic formulation

Estimate  $x$  from observation  $y \sim p(y|x)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

Data-fidelity

ex:  $f(x) = \frac{1}{2} \|Ax - y\|^2$

Regularization

ex: Total Variation [Rudin et al. '92]

Wavelet sparsity [Mallat '09]

Deep prior [Bora et al. '17]

## Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

Data-fidelity

ex:  $f(x) = \frac{1}{2} \|Ax - y\|^2$

Regularization

ex: Total Variation [Rudin et al. ‘92]

Wavelet sparsity [Mallat ‘09]

Deep prior [Bora et al. ‘17]

## Probabilistic formulation

Estimate  $x$  from observation  $y \sim p(y|x)$

**Maximum A-Posteriori**

$$x^* = \arg \max_{x \in \mathbb{R}^n} p(x|y)$$

$$= \arg \min_{x \in \mathbb{R}^n} -\log p(y|x) - \log p(x)$$

# Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

Data-fidelity

ex:  $f(x) = \frac{1}{2} \|Ax - y\|^2$

Regularization

ex: Total Variation [Rudin et al. ‘92]

Wavelet sparsity [Mallat ‘09]

Deep prior [Bora et al. ‘17]

# Probabilistic formulation

Estimate  $x$  from observation  $y \sim p(y|x)$

**Maximum A-Posteriori**

$$x^* = \arg \max_{x \in \mathbb{R}^n} p(x|y)$$

$$= \arg \min_{x \in \mathbb{R}^n} -\log p(y|x) - \log p(x)$$

- Gaussian noise:  $f(x) = \frac{1}{2} \|Ax - y\|^2$
- Poisson noise:  $f(x) = \text{KL}(y, Ax)$

# Variational formulation

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

💡 Estimation  $x^*$  must verify:

- $Ax^*$  “is close from”  $y$ .
- $x^*$  “looks like” a real image.

$$x^* = \arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

Data-fidelity

ex:  $f(x) = \frac{1}{2} \|Ax - y\|^2$

Regularization

ex: Total Variation [Rudin et al. '92]

Wavelet sparsity [Mallat '09]

Deep prior [Bora et al. '17]

# Probabilistic formulation

Estimate  $x$  from observation  $y \sim p(y|x)$

**Maximum A-Posteriori**

$$x^* = \arg \max_{x \in \mathbb{R}^n} p(x|y)$$

$$= \arg \min_{x \in \mathbb{R}^n} -\log p(y|x) - \log p(x)$$

- Gaussian noise:  $f(x) = \frac{1}{2} \|Ax - y\|^2$
- Poisson noise:  $f(x) = \text{KL}(y, Ax)$

$$g(x) = -\log p(x)$$



# Heuristics of Plug-and-Play methods [Venkatakrishnan et al. '13]

Estimate  $x$  from observation  $y \sim \text{noise}(Ax)$

- **Decouple** optimization over  $f$  and  $g$ . [Combette & Pesquet '11] [Zoran & Weiss '11]
- Use a **Gaussian Denoiser** as the optimization step over the regularizer  $g$ .
  - Easy Inverse problem  $A = \text{Id}$ , performant deep denoisers [Zhang et al. '17, '19, '21] [Song et al. '19]
  - $g(\text{Denoiser}(x)) < g(x)$  (**Implicit Regularization**)

From  $y \sim \text{noise}(Ax)$ , iterate:

1. Denoising step
2. Decrease the data-fidelity term  $f$ .

✓ “One denoiser to solve them all!”  
[Chang et. al]

# Presentation Outline

## 1. Introduction to Plug-and-Play methods

- ▶ How can an image Denoiser be used to regularize an inverse problem?

## 2. Convergence of Plug-and-Play algorithms

- ▶ Can sufficient conditions on the Denoiser be found to ensure that Plug-and-Play algorithms converge while maintaining performance?

## 3. Bregman Generalization and inverse problems with Poisson noise

- ▶ How can these schemes be applied to less regular data fidelity terms?

# Introduction to Plug-and-Play methods

# First-order optimization

Find  $\arg \min_{x \in \mathbb{R}^n} F(x)$

- Gradient Descent (**explicit**):

$$x_{k+1} = (\text{Id} - \tau \nabla F)(x_k)$$

- **Implicit** Gradient Descent /  
Proximal Point Algorithm:

$$\begin{aligned} x_{k+1} &= x_k - \tau \nabla F(x_{k+1}) \\ \Leftrightarrow x_{k+1} &= \text{Prox}_{\tau F}(x_k) \end{aligned}$$

where  $\text{Prox}_F(y) = \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|x - y\|^2 + F(x)$

# First-order optimization

Find  $\arg \min_{x \in \mathbb{R}^n} F(x)$

Find  $\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$

- Gradient Descent (explicit):

$$x_{k+1} = (\text{Id} - \tau \nabla F)(x_k)$$

- Implicit Gradient Descent /  
Proximal Point Algorithm:

$$x_{k+1} = x_k - \tau \nabla F(x_{k+1})$$

$$\Leftrightarrow x_{k+1} = \text{Prox}_{\tau F}(x_k)$$

where  $\text{Prox}_F(y) = \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|x - y\|^2 + F(x)$

- Gradient descent:

$$x_{k+1} = (\text{Id} - \tau \nabla(f + g))(x_k)$$

- Proximal Gradient Descent (PGD):

$$x_{k+1} = \text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)(x_k)$$

- Douglas-Rashford Splitting / ADMM:

$$x_{k+1} = \left( \frac{1}{2} \text{Id} + \frac{1}{2} (2\text{Prox}_{\tau g} - \text{Id}) \circ (2\text{Prox}_{\tau f} - \text{Id}) \right)$$

# First-order optimization

Find  $\arg \min_{x \in \mathbb{R}^n} F(x)$

Find  $\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$

- Gradient Descent (explicit):

$$x_{k+1} = (\text{Id} - \tau \nabla F)(x_k)$$

- Implicit Gradient Descent /  
Proximal Point Algorithm:

$$x_{k+1} = x_k - \tau \nabla F(x_{k+1})$$

$$\Leftrightarrow x_{k+1} = \text{Prox}_{\tau F}(x_k)$$

where  $\text{Prox}_F(y) = \arg \min_{x \in \mathbb{R}^n} \frac{1}{2} \|x - y\|^2 + F(x)$

- Gradient descent:

$$x_{k+1} = (\text{Id} - \tau \nabla(f + g))(x_k)$$

- Proximal Gradient Descent (PGD):

$$x_{k+1} = \text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)(x_k)$$

- Douglas-Rashford Splitting / ADMM:

$$x_{k+1} = \left( \frac{1}{2} \text{Id} + \frac{1}{2} (2\text{Prox}_{\tau g} - \text{Id}) \circ (2\text{Prox}_{\tau f} - \text{Id}) \right)$$

💡 Involve  $\nabla g$  or  $\text{Prox}_{\tau g}$  but not  $g$  explicitly.

# Denoising: an implicit regularization

Estimate  $x$  from observation  $y = x + \xi$

- $x \sim p(x)$ .
- $\xi \sim \mathcal{N}(0, \sigma^2 \text{Id})$  Gaussian noise .



# Denoising: an implicit regularization

Estimate  $x$  from observation  $y = x + \xi$

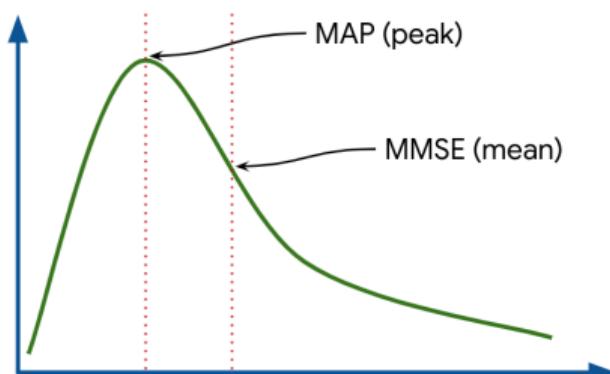
- $x \sim p(x)$ .
- $\xi \sim \mathcal{N}(0, \sigma^2 \text{Id})$  Gaussian noise .  
→ Two *theoretical* denoisers:

## MMSE denoiser

$$D_{\sigma}^{\text{MMSE}}(y) = \mathbb{E}_{x \sim p(x|y)}[x]$$

## MAP Denoiser

$$D_{\sigma}^{\text{MAP}}(y) = \arg \max_x p(x|y)$$



# Denoising: an implicit regularization

Estimate  $x$  from observation  $y = x + \xi$

- $x \sim p(x)$ .  $p_\sigma = p * G_\sigma$  with  $G_\sigma$  Gaussian pdf of variance  $\sigma^2$ .
- $\xi \sim \mathcal{N}(0, \sigma^2 \text{Id})$  Gaussian noise .  
→ Two *theoretical* denoisers:

## MMSE denoiser

$$D_\sigma^{\text{MMSE}}(y) = \mathbb{E}_{x \sim p(x|y)}[x]$$

## MAP Denoiser

$$D_\sigma^{\text{MAP}}(y) = \arg \max_x p(x|y)$$

Tweedie:

$$D_\sigma^{\text{MMSE}} = \text{Id} - \sigma^2 \nabla(-\log p_\sigma)$$

Bayes:

$$D_\sigma^{\text{MAP}} = \text{Prox}_{-\sigma^2 \log p}$$

# Denoising: an implicit regularization

Estimate  $x$  from observation  $y = x + \xi$

- $x \sim p(x)$ .  $p_\sigma = p * G_\sigma$  with  $G_\sigma$  Gaussian pdf of variance  $\sigma^2$ .
- $\xi \sim \mathcal{N}(0, \sigma^2 \text{Id})$  Gaussian noise .  
→ Two *theoretical* denoisers:

## MMSE denoiser

$$D_\sigma^{\text{MMSE}}(y) = \mathbb{E}_{x \sim p(x|y)}[x]$$

## MAP Denoiser

$$D_\sigma^{\text{MAP}}(y) = \arg \max_x p(x|y)$$

Tweedie:

$$D_\sigma^{\text{MMSE}} = \text{Id} - \sigma^2 \nabla(-\log p_\sigma)$$

Bayes:

$$D_\sigma^{\text{MAP}} = \text{Prox}_{-\sigma^2 \log p}$$

→ Given a *real* denoiser  $D_\sigma$ :

$$D_\sigma \approx D_\sigma^{\text{MMSE}} \quad \text{or} \quad D_\sigma \approx D_\sigma^{\text{MAP}}$$

Denoising ≈ explicit or implicit descent on  $g = -\log p$ .

## Small recap

1.  $g = -\log p$  is the ideal regularization
2. Minimize  $f + g$  requires access to  $\nabla g$  or  $\text{Prox } g$ .
3. A Gaussian denoiser approximates  $\nabla(-\log p)$  or  $\text{Prox}_{-\log p}$ .

## Plug-and-play algorithms

Estimate  $\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$  with  $f = -\log p(y|.)$  and  $g = -\log p$

$$\text{PGD: } \begin{cases} x_{k+1} = \text{Prox}_{\tau f} \circ (\text{Id} - \tau \nabla g)(x_k) \text{ or} \\ x_{k+1} = \text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)(x_k) \end{cases}$$

X unknown regularization  $g$

# Plug-and-play algorithms

Estimate  $\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$  with  $f = -\log p(y|.)$  and  $g = -\log p$

$$\text{PGD: } \begin{cases} x_{k+1} = \text{Prox}_{\tau f} \circ (\text{Id} - \tau \nabla g)(x_k) \text{ or} \\ x_{k+1} = \text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)(x_k) \end{cases}$$

✗ unknown regularization  $g$

## MMSE Denoiser

$$D_\sigma^{MMSE} = \text{Id} - \sigma^2 \nabla (-\log p_\sigma)$$

→ Given a *real* denoiser  $D_\sigma$ :

## GradPnP algorithms (RED)

[Romano et al., '16]

$$\nabla g \leftarrow \text{Id} - D_\sigma$$

## GradPnP-PGD:

$$x_{k+1} = \text{Prox}_{\tau f} \circ (\tau D_\sigma + (1 - \tau) \text{Id})(x_k)$$

# Plug-and-play algorithms

Estimate  $\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$  with  $f = -\log p(y|.)$  and  $g = -\log p$

$$\text{PGD: } \begin{cases} x_{k+1} = \text{Prox}_{\tau f} \circ (\text{Id} - \tau \nabla g)(x_k) \text{ or} \\ x_{k+1} = \text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)(x_k) \end{cases}$$

✗ unknown regularization  $g$

## MMSE Denoiser

$$D_\sigma^{MMSE} = \text{Id} - \sigma^2 \nabla (-\log p_\sigma)$$

## MAP Denoiser

$$D_\sigma^{MAP} = \text{Prox}_{-\sigma^2 \log p}$$

→ Given a *real* denoiser  $D_\sigma$ :

## GradPnP algorithms (RED)

[Romano et al., '16]

$$\nabla g \leftarrow \text{Id} - D_\sigma$$

## ProxPnP algorithms

[Venkatakrishnan et al., '13]

$$\text{Prox}_{\tau g} \leftarrow D_\sigma$$

### GradPnP-PGD:

$$x_{k+1} = \text{Prox}_{\tau f} \circ (\tau D_\sigma + (1 - \tau) \text{Id})(x_k)$$

### ProxPnP-PGD:

$$x_{k+1} = D_\sigma \circ (\text{Id} - \tau \nabla f)(x_k)$$

# Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

✓ Converges.

✗ Unknown regularization  $g$ .

## Plug-and-Play Algorithms

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

✓ Implicit regularization.

# Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

- ✓ Converges.
- ✗ Unknown regularization  $g$ .

## Plug-and-Play Algorithms

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.

# DPIR [Zhang et al. '21] Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

- ✓ Converges.
- ✗ Unknown regularization  $g$ .

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

## Plug-and-Play Algorithms

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.

$$\text{PSNR}(x_k, x^*)$$

Parameters  $\tau, \sigma$

$$\frac{\|x_{k+1} - x_k\|}{\|x_k\|}$$

# Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

✓ Converges.

✗ Unknown regularization  $g$ .

## Plug-and-Play Algorithms

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.
- ✗ No convergence guarantees.

✗ For a generic denoiser (e.g.: a deep neural net):

$$D_\sigma \neq \text{Id} - \nabla g \text{ and } D_\sigma \neq \text{Prox}_{\tau g}.$$

DPIR [Zhang et al. '21] + 500 iterations with fixed  $\tau$  and  $\sigma$ .  
Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

### MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

- ✓ Converges.
- ✗ Unknown regularization  $g$ .

### Plug-and-Play Algorithms

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.
- ✗ No convergence guarantees.

$$\text{PSNR}(x_k, x^*)$$

✗ For a generic denoiser (e.g.: a deep neural net):

$$D_\sigma \neq \text{Id} - \nabla g \text{ and } D_\sigma \neq \text{Prox}_{\tau g}.$$

Parameters  $\tau, \sigma$

$$\frac{\|x_{k+1} - x_k\|}{\|x_k\|}$$

# Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

- ✓ Converges.
- ✗ Unknown regularization  $g$ .

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

## Plug-and-Play Algorithms

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.
- ✗ No convergence guarantees.
- ✗ No minimized functional.

✗ For a generic denoiser (e.g.: a deep neural net):

$$D_\sigma \neq \text{Id} - \nabla g \text{ and } D_\sigma \neq \text{Prox}_{\tau g}.$$

# Plug-and-play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

- ✓ Converges.
- ✗ Unknown regularization  $g$ .

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

## Plug-and-Play Algorithms

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.
- ✗ No convergence guarantees.
- ✗ No minimized functional.

✗ For a generic denoiser (e.g.: a deep neural net):

$$D_\sigma \neq \text{Id} - \nabla g \text{ and } D_\sigma \neq \text{Prox}_{\tau g}.$$

**Objectif:** Find sufficient conditions on the denoiser to ensure convergence of the PnP algorithms.

## Existing works: fixed-point convergence

Show that PnP converges towards  $x^* \in \text{Fix}(T_{PnP})$ .

$$\begin{aligned}x_{k+1} &= \text{Prox}_{\tau f} \circ (\tau D_\sigma + (1 - \tau) \text{Id})(x_k) \\x_{k+1} &= D_\sigma \circ (\text{Id} - \tau \nabla f)(x_k)\end{aligned}$$

Show that  $T_{PnP}$  is non-expansive:

- $D_\sigma$  non-expansive (or averaged) [Sun et al. '19, Terris et al. '20, Hertrich et al. '21, Liu et al. '21, Bohra et al. '21]

$$\|D_\sigma(x) - D_\sigma(y)\| \leq \|x - y\|$$

✗ A performant generic deep denoiser is **not non-expansive**.

- data-fidelity term  $f$  strictly convex [Ryu et al., '19]  
✗ Excludes numerous problems

# PnP algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

✓ Converges.

✗ Unknown regularization  $g$ .

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

## Plug-and-Play algorithms

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.
- ✗ No convergence guarantees.
- ✗ No minimized functional.

$D_\sigma$  non-expansif

## Fixed-point convergence

$$x^* \in \text{Fix}(T_{PnP})$$

- ✓ Convergence.
- ✗ Sub-optimal performance .
- ✗ No minimized functional.

# Convergence via minimization

Prove that PnP algorithms converge towards  $x^* \in \arg \min_x f(x) + g(x)$ .

✗ For a generic denoiser (e.g.: a deep neural network):

$$D_\sigma \neq \text{Id} - \nabla g \text{ et } D_\sigma \neq \text{Prox}_g.$$

**Solution:** Make PnP real optimization algorithms by guaranteeing exactly

**GradPnP algorithms**

$$D_\sigma = \text{Id} - \nabla g$$

**ProxPnP algorithms**

$$D_\sigma = \text{Prox}_g$$

# Convergence via minimization

Prove that PnP algorithms converge towards  $x^* \in \arg \min_x f(x) + g(x)$ .

✗ For a generic denoiser (e.g.: a deep neural network):

$$D_\sigma \neq \text{Id} - \nabla g \text{ et } D_\sigma \neq \text{Prox}_g.$$

**Solution:** Make PnP real optimization algorithms by guaranteeing exactly

GradPnP algorithms

$$D_\sigma = \text{Id} - \nabla g$$

] ProxPnP algorithms

$$D_\sigma = \text{Prox}_g$$

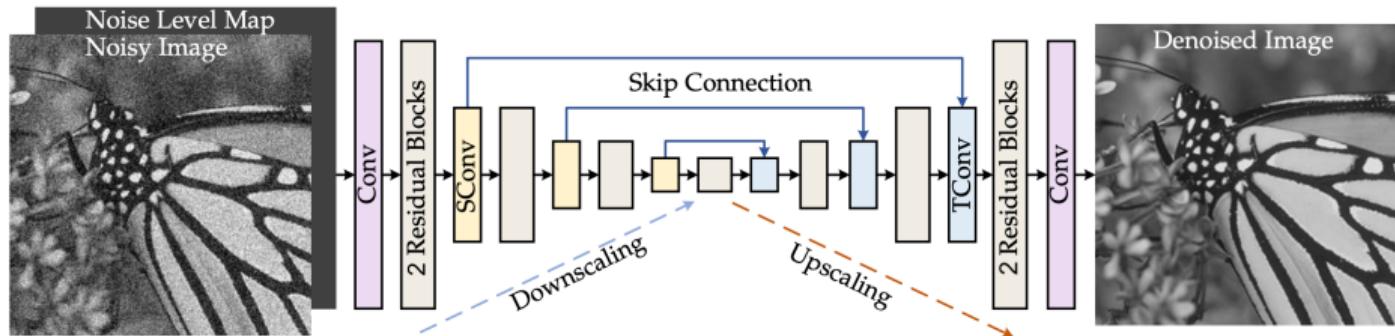
## Gradient-Step Denoiser [H., Leclaire, Papadakis, '21]

**Gradient-Step Denoiser (GS)**  $D_\sigma = \text{Id} - \nabla g_\sigma$

# Gradient-Step Denoiser [H., Leclaire, Papadakis, '21]

**Gradient-Step Denoiser (GS)**  $D_\sigma = \text{Id} - \nabla g_\sigma$

- With the **non-convex** potential:  $g_\sigma(x) = \frac{1}{2} \|x - N_\sigma(x)\|^2$   
 $N_\sigma : \mathbb{R}^n \rightarrow \mathbb{R}^n$  UNet with differentiable activations



[Zhang et al. '21]

$$D_\sigma(x) = x - \nabla \left( \frac{1}{2} \|x - N_\sigma(x)\|^2 \right) = N_\sigma(x) + J_{N_\sigma}(x)^T(x - N_\sigma(x))$$

# Gradient-Step Denoiser [H., Leclaire, Papadakis, '21]

**Gradient-Step Denoiser (GS)**  $D_\sigma = \text{Id} - \nabla g_\sigma$

- With the **non-convex** potential:  $g_\sigma(x) = \frac{1}{2}||x - N_\sigma(x)||^2$
- Trained by minimizing the  $L^2$  cost
- ✓ **No degradation** of denoising performance

| $\sigma(./255)$ | 15    | 25    | 50    |
|-----------------|-------|-------|-------|
| UNet            | 33.90 | 31.25 | 28.00 |
| GS-UNet         | 33.90 | 31.26 | 28.01 |

PSNR on the CBSD68 dataset

# Gradient-Step Denoiser [H., Leclaire, Papadakis, '21]

**Gradient-Step Denoiser (GS)**  $D_\sigma = \text{Id} - \nabla g_\sigma$

- With the **non-convex** potential:  $g_\sigma(x) = \frac{1}{2} \|x - N_\sigma(x)\|^2$
- Trained by minimizing the  $L^2$  cost
- ✓ **No degradation** of denoising performance

| $\sigma(./255)$ | 15    | 25    | 50    |
|-----------------|-------|-------|-------|
| UNet            | 33.90 | 31.25 | 28.00 |
| GS-UNet         | 33.90 | 31.26 | 28.01 |

PSNR on the CBSD68 dataset

- The MMSE denoiser is optimal for the  $L^2$  cost:

$$D_\sigma = \text{Id} - \nabla g_\sigma \approx D_\sigma^{MMSE} = \text{Id} - \nabla(-\sigma^2 \log p_\sigma)$$

✓  $g_\sigma \approx -\sigma^2 \log p_\sigma + \text{const}$

# Nonconvex convergence of GradPnP algorithms [H., Leclaire, Papadakis, '21]

With GS denoiser, GradPnP algorithms converge towards a critical point of  $f + \lambda g_\sigma$ .

ex: GradPnP-PGD

$$\begin{aligned}x_{k+1} &= \text{Prox}_{\tau f} \circ (\tau \lambda D_\sigma + (1 - \tau \lambda) \text{Id})(x_k) \\&= \text{Prox}_{\tau f} \circ (\text{Id} - \tau \lambda \nabla g_\sigma)(x_k)\end{aligned}$$

- ✗  $g_\sigma$  is **non-convex**.
- ✓  $g_\sigma$  has Lipschitz gradient

## Hypotheses

- (i)  $f$  convex, proper, l.s.c
- (ii)  $\nabla g_\sigma$   $L$ -Lipschitz ✓
- (iii)  $f + \lambda g_\sigma$  coercive and verifies the *Kurdyka-Łojasiewicz (KL)* property ✓

Then, for  $\tau < \frac{2\lambda}{L}$ ,  $(x_k)$  converges towards a critical point of  $f + \lambda g_\sigma$ . [Attouch et al., '13]

# Deblurring

(a)  $\frac{||x_{k+1} - x_k||}{||x_k||}$

(b)  $F(x_k) - F^*$

(c) PSNR( $x_k$ )

## Experimental results

| Input noise level         | Deblurring   |              |              | Super-resolution |              |              |
|---------------------------|--------------|--------------|--------------|------------------|--------------|--------------|
|                           | 0.01         | 0.03         | 0.05         | 0.01             | 0.03         | 0.05         |
| Bicubic                   | -            | -            | -            | 24.85            | 23.96        | 22.79        |
| EPLL [Zoran & Weiss, '11] | 27.34        | 25.16        | 24.04        | -                | -            | -            |
| IRCNN [Zhang et al., '17] | 31.42        | 28.01        | 26.40        | 26.97            | 25.86        | 25.45        |
| DPIR [Zhang et al., '21]  | <b>31.93</b> | <b>28.30</b> | <u>26.82</u> | 27.79            | 26.58        | <u>25.83</u> |
| GradPnP-PGD               | <u>31.70</u> | <u>28.28</u> | <b>26.86</b> | <u>27.88</u>     | <b>26.81</b> | <b>26.01</b> |

PSNR (dB) pour deblurring and super-resolution tasks, averaged over the CBSD68 dataset and over a variety of blur kernels.

- ✓ Convergence + State-of-the-art performance among Plug-and-Play methods.

# Convergence of PnP algorithms via minimization

Prouve that PnP algorithms converge towards  $x^* \in \arg \min_x f(x) + g(x)$ .

✗ For a generic denoiser (e.g.: a deep neural network):

$$D_\sigma \neq \text{Prox}_{\tau g}.$$

**Solution:** Make PnP algorithms real optimization algorithms by guaranteeing:

GradPnP algorithms

$$D_\sigma = \text{Id} - \nabla g$$

ProxPnP algorithms

$$D_\sigma = \text{Prox}_g$$

## Prox characterization

For an operator  $T : \mathbb{R}^n \rightarrow \mathbb{R}^n$ ,

[Moreau, '65]

$T = \nabla h$  with  $h$  convex and  $T$  1-Lipschitz  
 $\Leftrightarrow \exists \phi : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$  **convex** s.t.

$$T = \text{Prox}_\phi$$

$\times D_\sigma$  1-Lipschitz non-realistic.

## Prox characterization

For an operator  $T : \mathbb{R}^n \rightarrow \mathbb{R}^n$ ,

[Moreau, '65]

$T = \nabla h$  with  $h$  convex and  $T$  1-Lipschitz  
 $\Leftrightarrow \exists \phi : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$  **convex** s.t.

$$T = \text{Prox}_\phi$$

✗  $D_\sigma$  1-Lipschitz non-realistic.

[Gribonval et Nikolova, '20]

$T = \nabla h$  with  $h$  convex and  $T$  Lipschitz  
 $\Leftrightarrow \exists \phi : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$   
 $\left(1 - \frac{1}{\text{Lip}(T)}\right)$  **weakly convex** s.t.

$$T = \text{Prox}_\phi$$

✓ Verified by the MMSE denoiser

## Prox characterization

For an operator  $T : \mathbb{R}^n \rightarrow \mathbb{R}^n$ ,

[Moreau, '65]

$T = \nabla h$  with  $h$  convex and  $T$  1-Lipschitz  
 $\Leftrightarrow \exists \phi : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$  **convex** s.t.

$$T = \text{Prox}_\phi$$

✗  $D_\sigma$  1-Lipschitz non-realistic.

[Gribonval et Nikolova, '20]

$T = \nabla h$  with  $h$  convex and  $T$  Lipschitz  
 $\Leftrightarrow \exists \phi : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$   
 $\left(1 - \frac{1}{\text{Lip}(T)}\right)$  **weakly convex** s.t.

$$T = \text{Prox}_\phi$$

✓ Verified by the MMSE denoiser

Proximal Denoiser  
[H., Leclaire, Papadakis, '22]

**Gradient-Step Denoiser:**

$$D_\sigma = \text{Id} - \nabla g_\sigma = \nabla h_\sigma$$

$$\text{with } h_\sigma = \frac{1}{2} \|x\|^2 - g_\sigma$$

$\nabla g_\sigma$  1-Lipschitz  $\Rightarrow h_\sigma$  convex

Corollary

If  $\nabla g_\sigma = \text{Id} - D_\sigma$  is  $(L \leq 1)$ -Lipschitz  
 $\Rightarrow \exists \phi_\sigma : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$   
 $\frac{L}{L+1}$  weakly convex s.t.

$$D_\sigma = \text{Prox}_{\phi_\sigma}$$

$$\forall x \in \text{Im}(D_\sigma), \phi_\sigma = h_\sigma^*(x) - \frac{1}{2} \|x\|^2$$

# How to ensure $\text{Lip}(\nabla g_\sigma) \leq 1$ ? [H., Leclaire, Papadakis, '22]

**Proximal Denoiser**  $D_\sigma = \text{Id} - \nabla g_\sigma = \nabla h_\sigma = \text{Prox}_{\phi_\sigma}$

- Convex parametrization of  $h_\sigma$ . [Amos et al., '17] [Goujon et al., '23]  
✗ Suboptimal performance.
- Training cost penalization [Terris et. al, '21]:

$$\mathbb{E}_{x \sim p, y \sim \mathcal{N}(x, \sigma^2 \text{ Id})} \left[ \|D_\sigma(y) - x\|^2 + \mu \max(\|\nabla^2 g_\sigma(y)\|_S, 1 - \epsilon) \right]$$

| $\sigma(. / 255)$ | 5           | 15          | 25          |
|-------------------|-------------|-------------|-------------|
| GS-UNet           | 1.26        | 1.96        | 3.27        |
| <b>Prox-UNet</b>  | <b>0.92</b> | <b>0.99</b> | <b>0.96</b> |

$\max_x \|\nabla^2 g_\sigma(x + \xi_\sigma)\|_S$  on the CBSD68 dataset

| $\sigma(. / 255)$ | 5            | 15           | 25           |
|-------------------|--------------|--------------|--------------|
| GS-UNet           | <b>40.26</b> | <b>33.90</b> | <b>31.26</b> |
| <b>Prox-UNet</b>  | 40.12        | 33.60        | 30.82        |

PSNR on the CBSD68 dataset

# Convergence of ProxPnP algorithms [H., Leclaire, Papadakis, '22]

With Prox denoiser, ProxPnP algorithms converge towards a critical point of  $f + \lambda\phi_\sigma$ .

ex: ProxPnP-PGD

$$\begin{aligned}x_{k+1} &= D_\sigma \circ (\text{Id} - \frac{1}{\lambda} \nabla f)(x_k) \\&= \text{Prox}_{\phi_\sigma} \circ (\text{Id} - \frac{1}{\lambda} \nabla f)(x_k)\end{aligned}$$

⚠  $D_\sigma = \text{Prox}_{\phi_\sigma} \neq \text{Prox}_{\tau\phi_\sigma} \Rightarrow$  Requires to keep a stepsize  $\tau = 1$ .

## Hypothesis

- (i)  $f$  convex, differentiable, with  $L_f$ -Lipschitz gradient.
- (ii)  $\phi_\sigma$   $M$ -weakly convex. ✓
- (iii)  $f + \lambda\phi_\sigma$  coercive and verifies the *Kurdyka-Łojasiewicz (KL)* condition. ✓

Then, for  $\frac{1}{\lambda}L_f + M < 2$ ,  $(x_k)$  converges towards a critical point of  $f + \lambda\phi_\sigma$ .

# Super-resolution

(a)  $\frac{||x_{k+1} - x_k||}{||x_k||}$

(b)  $F(x_k) - F^*$

(c) PSNR( $x_k$ )

# Plug-and-Play algorithms

Estimate  $x$  from observation  $y \sim p(y|x)$

## MAP / 1<sup>st</sup> order optim.

$$\arg \min_{x \in \mathbb{R}^n} f(x) + g(x)$$

- ✓ Converging algorithms.
- ✗ Unknown regularization  $g$ .

$$\frac{D_\sigma \approx \text{Id} - \nabla g}{D_\sigma \approx \text{Prox}_{\tau g}}$$

## Plug-and-Play algorithms

- ✓ Implicit regularization.
- ✓ State-of-the-art performance.
- ✗ No convergence guarantees.
- ✗ No minimized functional.

## GradPnP convergence

$$x^* \in \arg \min_{x \in \mathbb{R}^n} f(x) + g_\sigma(x)$$

- ✓ Convergence.
- ✓ Optimize an explicit functional.
- ✓ State-of-the-art performance.

$$D_\sigma = \text{Id} - \nabla g_\sigma$$

$$D_\sigma = \text{Prox}_{\phi_\sigma}$$

## ProxPnP convergence

$$x^* \in \arg \min_{x \in \mathbb{R}^n} f(x) + \phi_\sigma(x)$$

$D_\sigma$  non-expansif

## Fixed-point convergence

$$x^* \in \text{Fix}(T_{PnP})$$

- ✓ Convergence.
- ✗ No minimized functional.
- ✗ Sub-optimal performance.

# Bregman generalization and inverse problems with Poisson noise.

[H., Kamilov, Leclaire, Papadakis, '23]

# Poisson inverse problems

If the observation noise is Poisson,  $y \sim \mathcal{P}(Ax)$ , the data-fidelity term is

$$f(x) = -\log p(y|x) = \text{KL}(y, Ax)$$

✗  $\nabla f$  non-Lipschitz.

[Bauschke et. al, '17] For Burgs-entropy  $h(x) = -\sum_i \log(x_i)$  and  $L \geq \|y\|_1$ :

✓ (L-smad)  $Lh - f$  convex on  $\text{int dom}(h)$ .

- Generalizes the notion of **smoothness**:

for  $h(x) = \frac{1}{2}\|x\|^2$ , implies  $\nabla f$   $L$ -Lipschitz.

- $h$  defines a **new geometry (Hessian manifold) adapted to the data-fidelity term  $f$** :

Euclidian distance  $\frac{1}{2}\|x - y\|^2$  replaced by the Bregman divergence  
 $D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$

*Bregman potential*

Euclidian  $\frac{1}{2}||x||^2$

$h : \mathbb{R}^n \rightarrow \mathbb{R}$  strictly convex,  $\mathcal{C}^2$

---

$${}^1D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$$

|                          |                                |  |
|--------------------------|--------------------------------|--|
| <i>Bregman potential</i> | Euclidian $\frac{1}{2}\ x\ ^2$ | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>          | $\frac{1}{2}\ x - y\ ^2$       | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>            | $\nabla f$ $L$ -Lipschitz      | $Lh - f$ convex  |

$${}^1D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$$

|                          |  |  |
|--------------------------|--|--|
| <i>Bregman potential</i> | Euclidian $\frac{1}{2}\ x\ ^2$   | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>          | $\frac{1}{2}\ x - y\ ^2$   | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>            | $\nabla f$ $L$ -Lipschitz  | $Lh - f$ convex  |
| <i>Proximal operator</i> | $\text{Prox}_f = \arg \min_x f(x) + \frac{1}{2}\ x - y\ ^2$ $\text{Prox}_f^h = \arg \min_x f(x) + D_h(x, y)$ |  |

$${}^1D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$$

|                                |   |  |
|--------------------------------|---|--|
| <i>Bregman potential</i>       | Euclidian $\frac{1}{2}\ x\ ^2$                              | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>                | $\frac{1}{2}\ x - y\ ^2$                                    | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>                  | $\nabla f$ $L$ -Lipschitz                                   | $Lh - f$ convex  |
| <i>Proximal operator</i>       | $\text{Prox}_f = \arg \min_x f(x) + \frac{1}{2}\ x - y\ ^2$ | $\text{Prox}_f^h = \arg \min_x f(x) + D_h(x, y)$                           |
| <i>Gradient Desc. (GD)</i>     | $\text{Id} - \tau(\nabla f + \nabla g)$                     | $\nabla h^*(\nabla h - \tau(\nabla f + \nabla g))$                         |
| <i>Prox. Grad. Desc. (PGD)</i> | $\text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)$    | $\text{Prox}_{\tau g}^h \circ \nabla h^*(\nabla h - \tau \nabla f)$        |

[H., Kamilov, Leclaire, Papadakis, '23]

$${}^1 D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$$

|   |   |  |
|---|---|--|
| <i>Bregman potential</i>                | Euclidian $\frac{1}{2}\ x\ ^2$                              | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>                         | $\frac{1}{2}\ x - y\ ^2$                                    | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>                           | $\nabla f$ $L$ -Lipschitz                                   | $Lh - f$ convex  |
| <i>Proximal operator</i>                | $\text{Prox}_f = \arg \min_x f(x) + \frac{1}{2}\ x - y\ ^2$ | $\text{Prox}_f^h = \arg \min_x f(x) + D_h(x, y)$                           |
| <i>Gradient Desc. (GD)</i>              | $\text{Id} - \tau(\nabla f + \nabla g)$                     | $\nabla h^*(\nabla h - \tau(\nabla f + \nabla g))$                         |
| <i>Prox. Grad. Desc. (PGD)</i>          | $\text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)$    | $\text{Prox}_{\tau g}^h \circ \nabla h^*(\nabla h - \tau \nabla f)$        |
| [H., Kamilov, Leclaire, Papadakis, '23] |   |  |
| <i>Noise model</i> $p(y x)$             | Gaussian  | $\propto \exp(-\gamma D_h(x, y))$  |

<sup>1</sup> $D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$

|                                |   |  |
|--------------------------------|---|--|
| <i>Bregman potential</i>       | Euclidian $\frac{1}{2}\ x\ ^2$                              | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>                | $\frac{1}{2}\ x - y\ ^2$                                    | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>                  | $\nabla f$ $L$ -Lipschitz                                   | $Lh - f$ convex  |
| <i>Proximal operator</i>       | $\text{Prox}_f = \arg \min_x f(x) + \frac{1}{2}\ x - y\ ^2$ | $\text{Prox}_f^h = \arg \min_x f(x) + D_h(x, y)$                           |
| <i>Gradient Desc. (GD)</i>     | $\text{Id} - \tau(\nabla f + \nabla g)$                     | $\nabla h^*(\nabla h - \tau(\nabla f + \nabla g))$                         |
| <i>Prox. Grad. Desc. (PGD)</i> | $\text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)$    | $\text{Prox}_{\tau g}^h \circ \nabla h^*(\nabla h - \tau \nabla f)$        |

[H., Kamilov, Leclaire, Papadakis, '23]

|                             |  |   |
|-----------------------------|--|---|
| <i>Noise model</i> $p(y x)$ | Gaussian                               | $\propto \exp(-\gamma D_h(x, y))$                                       |
| <i>MAP denoiser</i>         | $\text{Prox}_{-\sigma^2 \log p}(y)$    | $\text{Prox}_{-\frac{1}{\gamma} \log p}^h(y)$                           |
| <i>MMSE denoiser</i>        | $y + \sigma^2 \nabla \log p_\sigma(y)$ | $y + \frac{1}{\gamma} \nabla^2 h(y)^{-1} \cdot \nabla \log p_\gamma(y)$ |

<sup>1</sup> $D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$

|   |   |  |
|---|---|--|
| <i>Bregman potential</i>                | Euclidian $\frac{1}{2}\ x\ ^2$                              | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>                         | $\frac{1}{2}\ x - y\ ^2$                                    | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>                           | $\nabla f$ $L$ -Lipschitz                                   | $Lh - f$ convex  |
| <i>Proximal operator</i>                | $\text{Prox}_f = \arg \min_x f(x) + \frac{1}{2}\ x - y\ ^2$ | $\text{Prox}_f^h = \arg \min_x f(x) + D_h(x, y)$                           |
| <i>Gradient Desc. (GD)</i>              | $\text{Id} - \tau(\nabla f + \nabla g)$                     | $\nabla h^*(\nabla h - \tau(\nabla f + \nabla g))$                         |
| <i>Prox. Grad. Desc. (PGD)</i>          | $\text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)$    | $\text{Prox}_{\tau g}^h \circ \nabla h^*(\nabla h - \tau \nabla f)$        |
| [H., Kamilov, Leclaire, Papadakis, '23] |   |  |
| <i>Noise model</i> $p(y x)$             | Gaussian  | $\propto \exp(-\gamma D_h(x, y))$  |
| <i>MAP denoiser</i>                     | $\text{Prox}_{-\sigma^2 \log p}(y)$                         | $\text{Prox}_{-\frac{1}{\gamma} \log p}^h(y)$                              |
| <i>MMSE denoiser</i>                    | $y + \sigma^2 \nabla \log p_\sigma(y)$                      | $y + \frac{1}{\gamma} \nabla^2 h(y)^{-1} \cdot \nabla \log p_\gamma(y)$    |
| <i>GradPnP algorithms: GD with ...</i>  | $\nabla g \leftarrow \text{Id} - D_\sigma$                  | $\nabla g(y) \leftarrow \nabla^2 h(y) \cdot (\text{Id} - D_\gamma)(y)$     |
| <i>ProxPnP algorithms: PGD with ...</i> | $\text{Prox}_{\tau g} \leftarrow D_\sigma$                  | $\text{Prox}_{\tau g}^h \leftarrow D_\gamma$                               |

<sup>1</sup> $D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$

|                                |   |  |
|--------------------------------|---|--|
| <i>Bregman potential</i>       | Euclidian $\frac{1}{2}\ x\ ^2$                              | $h : \mathbb{R}^n \rightarrow \mathbb{R}$ strictly convex, $\mathcal{C}^2$ |
| <i>Distance</i>                | $\frac{1}{2}\ x - y\ ^2$                                    | Bregman Divergence $D_h(x, y)$ <sup>1</sup>                                |
| <i>Smooth</i>                  | $\nabla f$ $L$ -Lipschitz                                   | $Lh - f$ convex  |
| <i>Proximal operator</i>       | $\text{Prox}_f = \arg \min_x f(x) + \frac{1}{2}\ x - y\ ^2$ | $\text{Prox}_f^h = \arg \min_x f(x) + D_h(x, y)$                           |
| <i>Gradient Desc. (GD)</i>     | $\text{Id} - \tau(\nabla f + \nabla g)$                     | $\nabla h^*(\nabla h - \tau(\nabla f + \nabla g))$                         |
| <i>Prox. Grad. Desc. (PGD)</i> | $\text{Prox}_{\tau g} \circ (\text{Id} - \tau \nabla f)$    | $\text{Prox}_{\tau g}^h \circ \nabla h^*(\nabla h - \tau \nabla f)$        |

[H., Kamilov, Leclaire, Papadakis, '23]

|   |  |   |
|---|--|---|
| <i>Noise model</i> $p(y x)$             | Gaussian                                   | $\propto \exp(-\gamma D_h(x, y))$                                       |
| <i>MAP denoiser</i>                     | $\text{Prox}_{-\sigma^2 \log p}(y)$        | $\text{Prox}_{-\frac{1}{\gamma} \log p}^h(y)$                           |
| <i>MMSE denoiser</i>                    | $y + \sigma^2 \nabla \log p_\sigma(y)$     | $y + \frac{1}{\gamma} \nabla^2 h(y)^{-1} \cdot \nabla \log p_\gamma(y)$ |
| <i>GradPnP algorithms: GD with ...</i>  | $\nabla g \leftarrow \text{Id} - D_\sigma$ | $\nabla g(y) \leftarrow \nabla^2 h(y) \cdot (\text{Id} - D_\gamma)(y)$  |
| <i>ProxPnP algorithms: PGD with ...</i> | $\text{Prox}_{\tau g} \leftarrow D_\sigma$ | $\text{Prox}_{\tau g}^h \leftarrow D_\gamma$                            |

|                                  |  |   |
|----------------------------------|--|---|
| <i>Gradient-Step Denoiser</i>    | $D_\sigma(y) = y - \nabla g_\sigma(y)$ | $D_\gamma(y) = y - \nabla^2 h(y)^{-1} \cdot \nabla g_\gamma(y)$ |
| <i>Writes like a Prox if ...</i> | $h_\sigma$ convex                      | $J_{D_\gamma}$ positive semi-definite                           |

$$^1 D_h(x, y) = h(x) - h(y) - \langle \nabla h(y), x - y \rangle$$

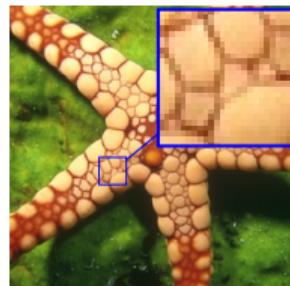
# Application to Poisson inverse problems

$$\text{Data-fidelity term } f(x) = KL(y, Ax)$$

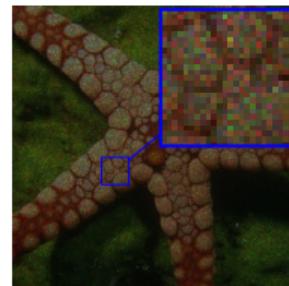
- **Bregman potential:** Burg's entropy  $h(x) = -\sum_{i=1}^n \log(x_i)$   
 $Lh - f$  convex for  $L \geq \|y\|_1$  [Bauschke, '17]
- **Bregman noise model:** Gamma Inverse ( $\mathcal{IG}$ ) distribution

$$p(y|x) \propto \prod_{i=1}^n \left( \frac{x_i}{y_i} \right)^\gamma \exp \left( -\gamma \frac{x_i}{y_i} \right) = \prod_{i=1}^n \mathcal{IG}(\gamma - 1, \gamma x_i)$$

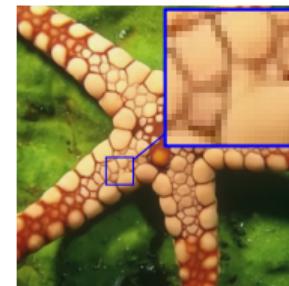
- **Gradient-Step Denoiser:**  $D_\gamma(y) = y - y^2 \nabla g_\gamma(y)$



(a) Proper

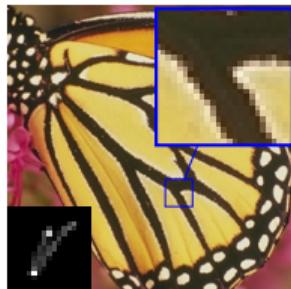


(b) noisy ( $\gamma = 25$ )

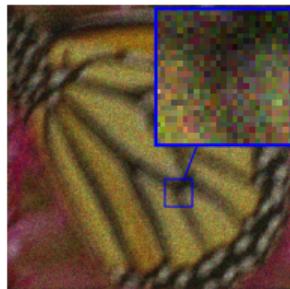


(c) Denoiser GS

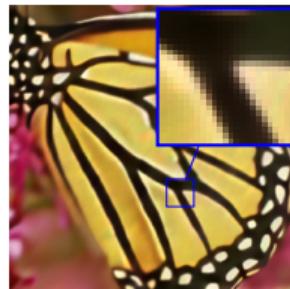
# Deblurring with Poisson noise



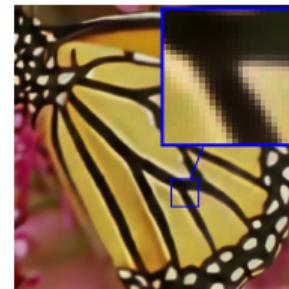
(a) Propre



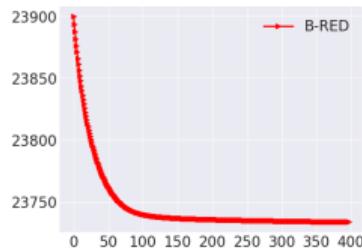
(b) Observation  
(16.21dB)



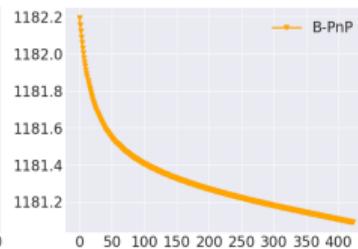
(c) GradPnP  
(23.01dB)



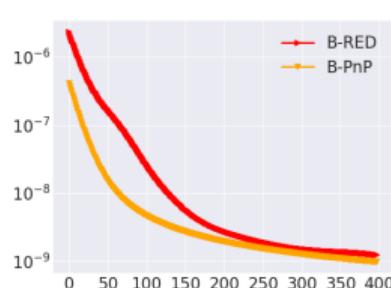
(d) ProxPnP  
(22.96dB)



(e)  $(\lambda f + g_\gamma)(x_k)$   
GradPnP



(f)  $(\lambda f + \phi_\gamma)(x_k)$   
ProxPnP



(g)  $\|x_{i+1} - x_i\|^2$

# Conclusion

## Summary of the contributions

- Introduction of denoisers allowing to rewrite PnP algorithms like real optimization algorithms.
- Learning of an explicit approximation of the clean image prior.
- Convergence proofs towards critical points of an explicit and non-convex objective.
- State-of-the-art restoration performance.
- Bregman extension for non-Gaussian noises.

# Open research questions

- **Convergence proofs with  $D_\sigma$  monotone but  $D_\sigma \neq$  gradient**

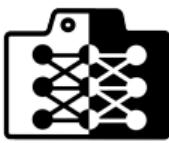
Convergence of the PnP algorithm  $D_\sigma \circ (\text{Id} - \nabla f)$  towards the solution of a non-monotone inclusion problem.

- **How to relate the error when training a denoiser to the performance in Plug-and-Play / Diffusion models ?**

If the trained denoiser is  $\epsilon$ -close to the optimal denoiser, do we still converge ? And if yes, to which point / distribution ?

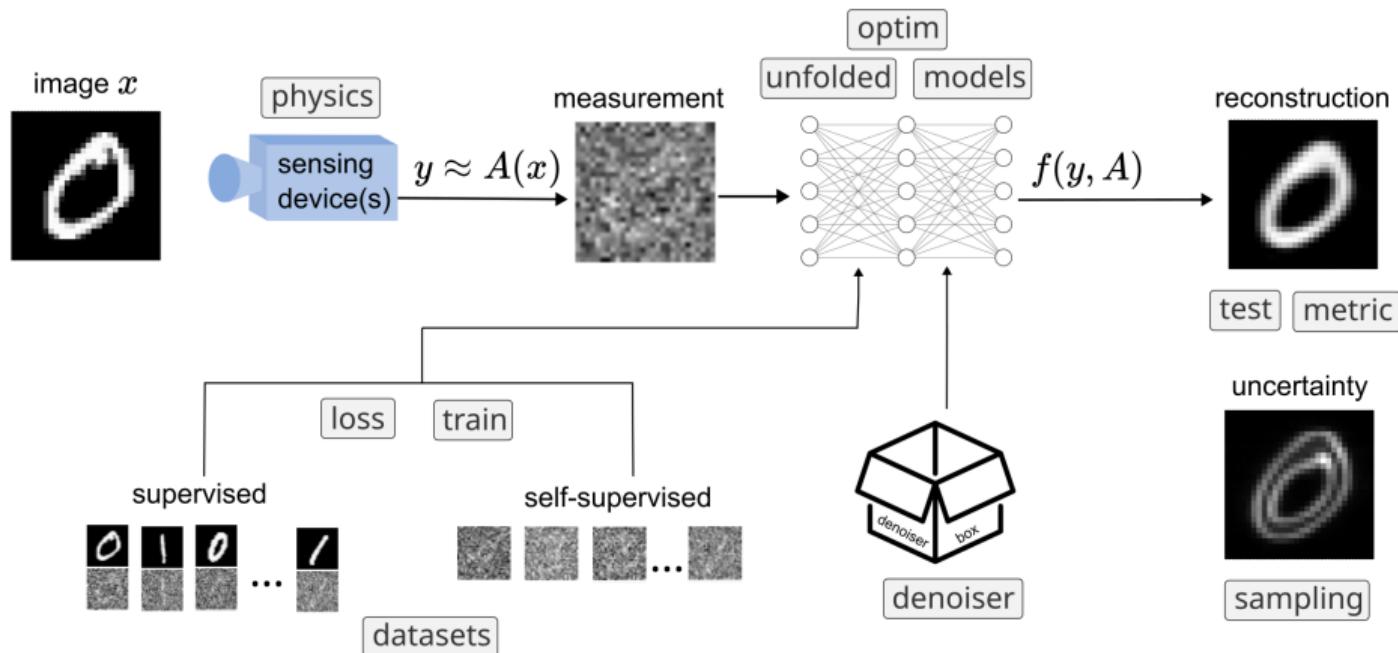
- **Denoising and Optimal Transport**

The MMSE denoiser (optimal denoiser for the  $L_2$  cost) writes as  $D_\sigma = \nabla m_\sigma$  with  $m_\sigma$  convex. *Brenier theorem:* It is an optimal transport map from  $p_\sigma = p * \mathcal{N}(0, \sigma^2 \text{Id})$  to  $D_\sigma \# p_\sigma$ . How to relate  $D_\sigma \# p_\sigma$  and  $p$  ?



# Deep Inverse

```
pip install deepinv
```



# Convergence proof with monotone denoiser

**Objective:** Generalize the nonconvex proof without requiring  $D_\sigma$  to write as a gradient.

| Optimization  | Monotone inclusion                                       |
|---|--|
| Find $\arg \min_x f(x) + g(x)$  | Find $x, 0 \in A(x) + B(x)$                              |
| PGD : $x_{k+1} \in \text{Prox}_{\tau f} \circ (\text{Id} - \tau \nabla g)(x_k)$ | $x_{k+1} \in J_{\tau A} \circ (\text{Id} - \tau B)(x_k)$ |
| $T = \partial f$ and $f$ convex   | $T$ (maximally) monotone                                 |
| $T = \nabla f$ and $f$ convex and $T$ $L$ -Lipschitz                            | $T$ ( $\beta = \frac{1}{L}$ )-cocoercive                 |
| $T = \text{Prox}_f$ $f$ convex  | $T = J_A = (\text{Id} + A)^{-1}$ , $A$ monotone          |
| $T = \partial f$ and $f$ weakly convex  | $T$ negative-monotone                                    |
| $T = \text{Prox}_f$ and $f$ weakly convex                                       | $T = J_A$ , $A$ negative-monotone                        |

$T$  monotone :  $\langle x - y, T(x) - T(y) \rangle \geq 0$

$T$  cocoercive :  $\langle x - y, T(x) - T(y) \rangle \geq \rho \|T(x) - T(y)\|^2$

$T$  negative monotone :  $\langle x - y, T(x) - T(y) \rangle \geq -\rho \|x - y\|^2$  ( $\rho \geq 0$ )

[Gribonval et Nikolova, '20]

$T = \nabla h$  with  $h$  convex and  $T$  Lipschitz  
 $\Leftrightarrow \exists \phi : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\} \left( 1 - \frac{1}{\text{Lip}(T)} \right)$   
**weakly convex** s.t.

$$T = \text{Prox}_\phi$$

[Bauschke et. al , '21]

$T$   $\beta$ -cocoercive  
 $\Leftrightarrow \exists A : \mathbb{R}^n \rightarrow \mathbb{R}^n \ (\rho = 1 - \beta)$   
**negative-monotone** s.t.

$$T = J_A = (\text{Id} + A)^{-1}$$

$$T \text{ } \beta\text{-cocoercive} \Leftrightarrow \frac{2}{\beta}T - \text{Id} \text{ } 1\text{-Lipschitz}$$

Objectives : Convergence of the PnP algorithm  $D_\sigma \circ (\text{Id} - \nabla f)$  towards the solution of a non-monotone inclusion problem.

- Study the cocoercive coefficient of state-of-the-art denoiser  $\Rightarrow$  resolvents of negative monotone operators.
- Convergence proof of the forward-backward algorithm for non-monotone inclusion problems ?

# Convergence of ProxPnP algorithms [H., Leclaire, Papadakis, '22]

With Prox denoiser, ProxPnP algorithms converge towards a critical point of  $f + \lambda\phi_\sigma$ .

ex: ProxPnP-PGD

$$x_{k+1} = D_\sigma \circ (\text{Id} - \frac{1}{\lambda} \nabla f)(x_k)$$

$$= \text{Prox}_{\phi_\sigma} \circ (\text{Id} - \frac{1}{\lambda} \nabla f)(x_k)$$

⚠  $D_\sigma = \text{Prox}_{\phi_\sigma} \neq \text{Prox}_{\tau\phi_\sigma} \Rightarrow$  Requires to keep a stepsize  $\tau = 1$ .

## Hypothesis

- (i)  $f$  convex, differentiable, with  $L_f$ -Lipschitz gradient.
- (ii)  $\phi_\sigma$   $M$ -weakly convex. ✓
- (iii)  $f + \lambda\phi_\sigma$  coercive and verifies the *Kurdyka-Łojasiewicz (KL)* condition. ✓

Then, for  $\frac{1}{\lambda}L_f + M < 2$ ,  $(x_k)$  converges towards a critical point of  $f + \lambda\phi_\sigma$ .

# Convergence of ProxPnP algorithms [H., Leclaire, Papadakis, '22]

With Prox denoiser, ProxPnP algorithms converge towards a critical point of  $f + \lambda\phi_\sigma$ .

ex: ProxPnP-PGD

$$x_{k+1} = D_\sigma \circ (\text{Id} - \frac{1}{\lambda} \nabla f)(x_k)$$

$$= \text{Prox}_{\phi_\sigma} \circ (\text{Id} - \frac{1}{\lambda} \nabla f)(x_k)$$

⚠  $D_\sigma = \text{Prox}_{\phi_\sigma} \neq \text{Prox}_{\tau\phi_\sigma} \Rightarrow$  Requires to keep a stepsize  $\tau = 1$ .

## Hypothesis

- (i)  $f$  convex, differentiable, with  $L_f$ -Lipschitz gradient.
- (ii)  $\phi_\sigma$   $M$ -weakly convex. ✓
- (iii)  $f + \lambda\phi_\sigma$  coercive and verifies the *Kurdyka-Łojasiewicz (KL)* condition. ✓

Then, for  $\boxed{\frac{1}{\lambda}L_f + M < 2}$   $(x_k)$  converges towards a critical point of  $f + \lambda\phi_\sigma$ .

✗ Regularisation parameter  $\lambda$  lower-bounded: **suboptimal**.

## Relaxed PnP- $\alpha$ PGD algorithm [H., Chambolle, Leclaire, Papadakis, '23]

**Solution:** Inertial relaxation of the PGD algorithm ( $\alpha$ PGD):

$$\text{Pour } 0 < \alpha < 1 \quad \begin{cases} q_{k+1} &= (1 - \alpha)y_k + \alpha x_k \\ x_{k+1} &= D_\sigma^\gamma \circ (x_k - \frac{1}{\lambda} \nabla f(q_{k+1})) \\ y_{k+1} &= (1 - \alpha)y_k + \alpha x_{k+1}. \end{cases}$$

with proximal denoiser  $\gamma$ -relaxed  $D_\sigma^\gamma = \text{Id} - \gamma \nabla g_\sigma = \text{Prox}_{\phi_\sigma^\gamma}$

✓  $\phi_\sigma^\gamma$  is  $\frac{\gamma L}{\gamma L + 1}$  weakly convex: **control** of the weak convexity constant  $\gamma$ .

### Hypothesis

- (i)  $f$  convex, differentiable, with gradient  $L_f$ -Lipschitz.
- (ii)  $\phi_\sigma^\gamma$   $M_\gamma$ -weakly convex. ✓

Then, for  $\frac{1}{\lambda} L_f M_\gamma < 1$  and  $M_\gamma < \alpha < \frac{\lambda}{L_f}$ , ✓ Allow  $\lambda$  small when  $\gamma \rightarrow 0$

$(x_k)$  converges towards a critical point of  $f + \lambda \phi_\sigma^\gamma$ .

# Kurdyka-Łojasiewicz property

Definition (Kurdyka-Łojasiewicz (KL) property [Attouch et. al '10])

A function  $f : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{+\infty\}$  is said to have the Kurdyka-Łojasiewicz property at  $x^* \in \text{dom}(f)$  if there exists  $\eta \in (0, +\infty)$ , a neighborhood  $U$  of  $x^*$  and a continuous concave function  $\phi : [0, \eta) \rightarrow \mathbb{R}_+$  such that  $\phi(0) = 0$ ,  $\phi$  is  $\mathcal{C}^1$  on  $(0, \eta)$ ,  $\phi' > 0$  on  $(0, \eta)$  and  $\forall x \in U \cap [f(x^*) < f < f(x^*) + \eta]$ , the Kurdyka-Łojasiewicz inequality holds:

$$\phi'(f(x) - f(x^*)) \text{dist}(0, \partial f(x)) \geq 1. \quad (1)$$

Proper lsc functions that satisfy the Kurdyka-Łojasiewicz inequality at each point of  $\text{dom}(\partial f)$  are called KL functions.

For  $f$  smooth and  $f(x^*) = 0$ , the KL inequality equation 1 writes  $\|\nabla(\phi \circ f)(x)\| \geq 1$   
Useful set of functions with stability properties which verify the KL property:

- Real analytic functions
- Semialgebraic functions
- Subanalytic functions [Bolte et. al, '07]

## $\text{Id} - D_\sigma$ vs $D_\sigma$ non-expansif

| $\sigma(./255)$ | 5     | 10    | 15    | 20    | 25    |
|-----------------|-------|-------|-------|-------|-------|
| GS-DRUNet       | 40.27 | 36.16 | 33.92 | 32.41 | 31.28 |
| prox-DRUNet     | 40.04 | 35.86 | 33.51 | 31.88 | 30.64 |
| nonexp-DRUNet   | 34.92 | 32.90 | 31.42 | 30.30 | 29.42 |

PSNR (dB) de Débruitage on the CBSD68 dataset, pour différents niveaux de noises  $\sigma$ .

# Performance numérique des PnP algorithms avec Denoiser proximal

| $\sigma(./255)$       | 2.55  | 7.65  | 12.75 |
|-----------------------|-------|-------|-------|
| IRCNN                 | 31.42 | 28.01 | 26.40 |
| DPIR                  | 31.93 | 28.30 | 26.82 |
| GSRED-PGD             | 31.70 | 28.28 | 26.86 |
| ProxPnP-PGD           | 30.91 | 27.97 | 26.66 |
| ProxPnP- $\alpha$ PGD | 31.55 | 28.03 | 26.66 |
| ProxPnP-DRSdiff       | 30.56 | 27.78 | 26.61 |
| ProxPnP-DRS           | 31.51 | 28.01 | 26.63 |
| nonexp-PnP-PGD        | 30.25 | 27.06 | 25.30 |

PSNR (dB) lors du defloutage on the CBSD68 dataset, moyennés pour différents noyaux de flous

## Propriétés de $g_\sigma$

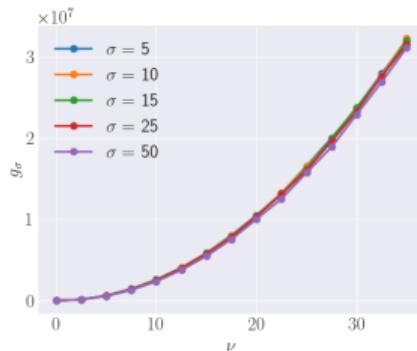
- $g_\sigma$  coercive:

On choisir  $C$  compact convexe ( $C = [-1, 2]^n$ ) et

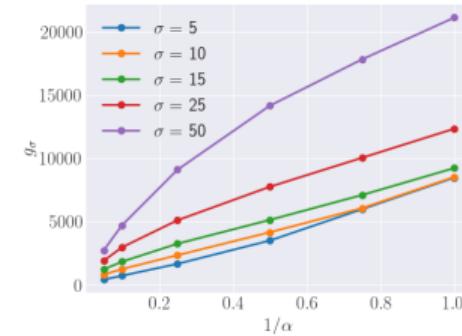
$$\hat{g}_\sigma(x) = g_\sigma(x) + \frac{1}{2} \|x - \Pi_C(x)\|^2 = \frac{1}{2} \|x - N_\sigma(x)\|^2 + \frac{1}{2} \|x - \Pi_C(x)\|^2$$

où  $\Pi_C$  la projection euclidienne sur  $C$ .

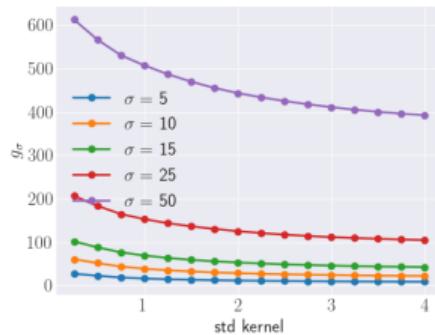
- $g_\sigma$  sous-analytique



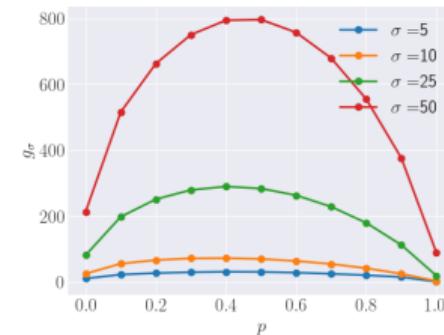
(a) Gaussian noise  $y \sim \mathcal{N}(x, \nu^2 \text{ Id})$



(b) Poisson noise  $y \sim \frac{1}{\alpha} \mathcal{P}(\alpha x)$



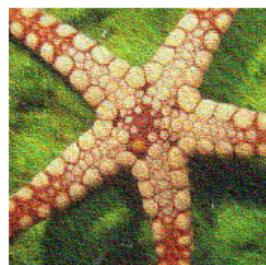
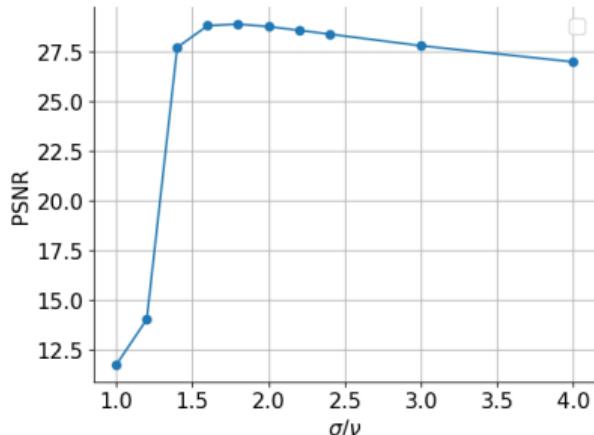
(c) Gaussian blur  $y = k_\nu * x$



(d) Missing pixels  $y = M_p x$

Evolution, for various  $\sigma$ , of  $g_\sigma(y)$  where  $y$  has been obtained from  $x$  with different degradations.

# Influence of the parameter $\sigma$



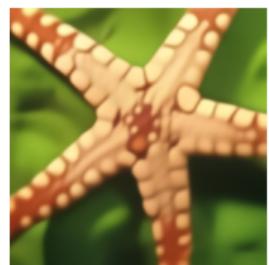
$\sigma/\nu = 1$



$\sigma/\nu = 1.8$

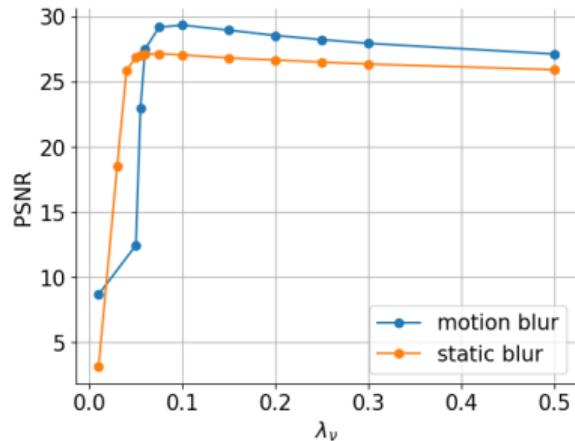


$\sigma/\nu = 4$



$\sigma/\nu = 10$

# Influence of the parameter $\lambda$



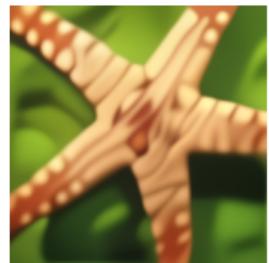
$$\lambda\nu^2 = 1$$



$$\lambda\nu^2 = 1.8$$



$$\lambda\nu^2 = 4$$



$$\lambda\nu^2 = 10$$

# Fixed-points of the GS-Denoiser

$$x_{k+1} = \mathcal{P}_{[0,1]^n}(x_k - \tau \nabla g_\sigma(x_k))$$



(a) init



(b)  $\sigma = 0$



(c)  $\sigma = 0.1$



(d)  $\sigma = 0.5$



(e)  $\sigma = 1$



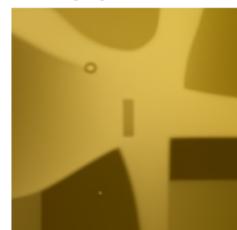
(f)  $\sigma = 5$



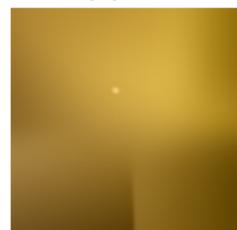
(g)  $\sigma = 10$



(h)  $\sigma = 20$



(i)  $\sigma = 30$



(j)  $\sigma = 40$

# Débruitage avec $N_\sigma$

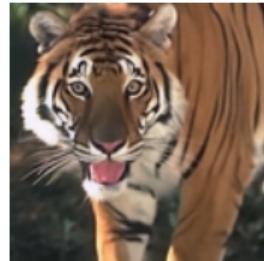
$$D_\sigma(x) = x - \nabla \left( \frac{1}{2} \|x - N_\sigma(x)\|^2 \right) = N_\sigma(x) + J_{N_\sigma}(x)^T(x - N_\sigma(x))$$



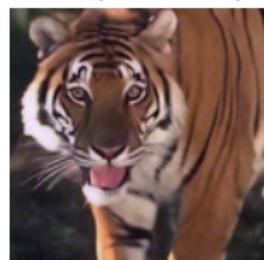
$\sigma = 10$



$\sigma = 25$



$N_\sigma$  (29.16dB)



$N_\sigma$  (23.94dB)



$D_\sigma$  (35.55dB)



$D_\sigma$  (30.49dB)

Denoiser relâché  $D_\sigma^\gamma = \text{Id} - \gamma \nabla g_\sigma$

Le Denoiser is entraîné pour que  $\nabla g_\sigma(x + \sigma n) \approx \sigma n$  où  $n \sim \mathcal{N}(0, \text{Id})$ .

$$D_\sigma^\gamma(x + \sigma n) \approx x + (1 - \gamma)\sigma n$$



$$\gamma = 0$$



$$\gamma = 0.1$$

A moderately sharp tiger image.

$$\gamma = 0.25$$



$$\gamma = 0.5$$



$$\gamma = 0.75$$

A perfectly sharp tiger image.

$$\gamma = 1$$

# Proximal Denoiser [Hurault, Leclaire, Papadakis, '22]

$$\text{Proximal Denoiser } D_\sigma = \text{Id} - \nabla g_\sigma = \text{Prox}_{\phi_\sigma}$$

$\nabla g_\sigma$   $L < 1$  Lipschitz  $\Rightarrow D_\sigma$  invertible

From  $y = D_\sigma(x)$ , one can recover  $x$  via fixed-point iterations

$$x_{k+1} = \nabla g_\sigma(x_k) + y$$



$x$



$D_\sigma(x)$



$x_\infty$

## PnP fixed-point convergence

$$x_{k+1} = T_{PnP}(x_k)$$

$$\text{with } T_{PnP} = \begin{cases} T_{HQS} &= D_\sigma \circ \text{Prox}_{\tau f} \\ T_{PGD} &= D_\sigma \circ (\text{Id} - \tau \nabla f) \\ T_{DRS} &= \frac{1}{2} \text{Id} + \frac{1}{2} (2D_\sigma - \text{Id}) \circ (2 \text{Prox}_{\tau f} - \text{Id}) \end{cases}$$

**Objective:** Show that  $x_k \rightarrow x^* \in \text{Fix}(T_{PnP})$ .

**Solution:** Averaged (nonexpansive)  $D_\sigma$ . [Sun et. al '18, '19, '21] [Hertrich et. al '20] [Terris et. al '21] [Bohra et. al '21]

- For  $\theta \in (0, 1)$ ,  $T : \mathbb{R}^n \rightarrow \mathbb{R}^n$   $\theta$ -averaged if  $T = \theta R + (1 - \theta) \text{Id}$  for some  $R$  nonexpansive (1-Lipschitz).
- If  $T$  is a  $\theta$ -averaged operator that admits fixed points, then the sequence  $x_{k+1} = T(x_k)$  converges to a fixed point of  $T$
- If  $D_\sigma$  is averaged then  $T_{PnP}$  is averaged for all presented PnP operators.



Non-expansivity harms denoising performance.

# Self-Consistent Velocity Matching of Probability Flows [Li, H., Solomon '23]

- We take PDEs of the form

$$\frac{\partial \mu_t}{\partial t}(x) = -\operatorname{div}(\mu_t f_t(x, \mu_t)) \quad (2)$$

where  $\mu_t$  is a probability density function. The Wasserstein Gradient Flow corresponds to  $f_t(x, \mu) = -\nabla_{W_2} D(\mu_t, \pi)(x)$ .

- $\mu_t$  and its velocity field  $v_t$  are linked via the **continuity equation**

$$\frac{\partial \mu_t}{\partial t} = -\operatorname{div}(\mu_t v_t). \quad (3)$$

- If a probability flow  $\mu_t$  with velocity field  $v_t$  satisfies the following fixed-point equation

$$v_t(x) = f_t(x, \mu_t) \quad (4)$$

then the PDE is solved.

We parameterize  $v_t^\theta$  with a time-dependent neural network and minimize

$$\min_\theta \int_0^T \mathbb{E}_{X \sim \mu_t^\theta} [| | v_t^\theta(X) - f_t(X, \mu_t^\theta) | |^2] dt. \quad (5)$$